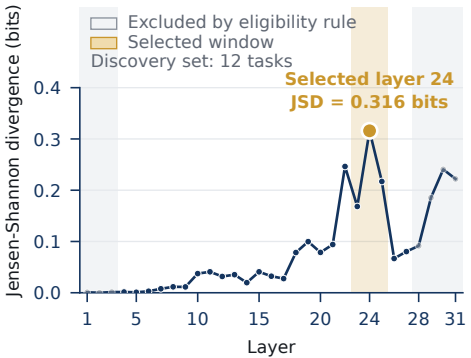
(a) Layerwise Jacobian-Lens readout divergence
Excluded by eligibility rule
Selected window
Discovery set: 12 tasks
Selected layer 24
JSD = 0.316 bits
Jensen-Shannon divergence (bits)
0.0
0.1
0.2
0.3
0.4
1
5
10
15
20
24
28
31
Layer

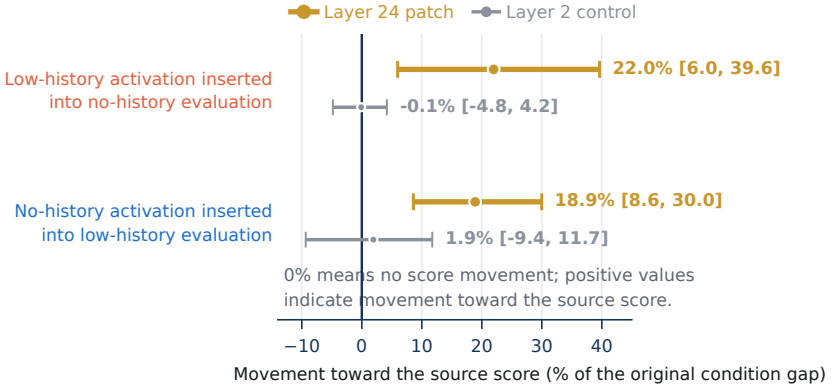
(b) Layer 24 transfers about one-fifth of the score gap
Layer 24 patch
Layer 2 control
Low-history activation inserted into no-history evaluation
22.0% [6.0, 39.6]
-0.1% [-4.8, 4.2]
No-history activation inserted into low-history evaluation
18.9% [8.6, 30.0]
1.9% [-9.4, 11.7]
0% means no score movement; positive values indicate movement toward the source score.
−10
0
10
20
30
40
Movement toward the source score (% of the original condition gap)